

# LIKWID: A lightweight performance-oriented tool suite for x86 multicore environments

Jan Treibig, Georg Hager, Gerhard Wellein  
 Erlangen Regional Computing Center (RRZE)  
 University of Erlangen-Nuremberg  
 Erlangen, Germany  
 Email: jan.treibig@rrze.uni-erlangen.de

**Abstract**—Exploiting the performance of today’s processors requires intimate knowledge of the microarchitecture as well as an awareness of the ever-growing complexity in thread and cache topology. LIKWID is a set of command-line utilities that addresses four key problems: Probing the thread and cache topology of a shared-memory node, enforcing thread-core affinity on a program, measuring performance counter metrics, and toggling hardware prefetchers. An API for using the performance counting features from user code is also included. We clearly state the differences to the widely used PAPI interface. To demonstrate the capabilities of the tool set we show the influence of thread pinning on performance using the well-known OpenMP STREAM triad benchmark, and use the affinity and hardware counter tools to study the performance of a stencil code specifically optimized to utilize shared caches on multicore chips.

## I. INTRODUCTION

Today’s multicore x86 processors bear multiple complexities when aiming for high performance. Conventional performance tuning tools like Intel VTune, OProfile, CodeAnalyst, OpenSpeedshop, etc., require a lot of experience in order to get sensible results. For this reason they are usually unsuitable for the scientific user, who would often be satisfied with a rough overview of the performance properties of their application code. Moreover, advanced tools often require kernel patches and additional software components, which makes them unwieldy and bug-prone. Additional confusion arises with the complex multicore, multicache, multisocket structure of modern systems (see Fig. 1); users are all too often at a loss about how hardware thread IDs are assigned to resources like cores, caches, sockets and NUMA domains. Moreover, the technical details of how threads and processes are bound to those resources vary strongly across compilers and MPI libraries.

LIKWID (“Like I Knew What I’m Doing”) is a set of easy to use command line tools to support optimization. It is targeted towards performance-oriented programming in a Linux environment, does not require any kernel patching, and is suitable for Intel and AMD processor architectures. Multithreaded and even hybrid shared/distributed-memory parallel code is supported. It comprises the following tools:

- `likwid-features` can display and alter the state of the on-chip hardware prefetching units in Intel x86 processors.
- `likwid-topology` probes the hardware thread and cache topology in multicore, multisocket nodes. Knowledge like this is required to optimize resource usage like, e.g., shared caches and data paths, physical cores, and ccNUMA locality domains, in parallel code.
- `likwid-perfCtr` measures performance counter metrics over the complete runtime of an application or, with support from a simple API, between arbitrary points in the code. Counter multiplexing allows the concurrent measurement of a large number of metrics, larger than the (usually small) number of available counters. Although it is possible to specify the full, hardware-dependent event names, some predefined event sets simplify matters when standard information like memory bandwidth or Flop counts is needed.
- `likwid-pin` enforces thread-core affinity in a multithreaded application “from the outside,” i.e., without changing the source code. It works with all threading models that are based on POSIX threads, and is also compatible with hybrid “MPI+threads” programming. Sensible use of `likwid-pin` requires correct information about thread numbering and cache topology, which can be delivered by `likwid-topology` (see above).

Although the four tools may appear to be partly unrelated, they solve the typical problems application programmers have when porting and running their code on complex multicore/multisocket environments. Hence, we consider it a natural idea to provide them as a single tool set.

This paper is organized as follows. Section II describes the four tools in some detail and gives hints for typical use. In Section III we briefly compare LIKWID to the PAPI feature set. Section IV demonstrates the use of LIKWID in three different case studies, and Section V gives a summary and an outlook to future work.

## II. TOOLS

LIKWID only supports x86-based processors. Given the strong prevalence of those architectures in the HPC market

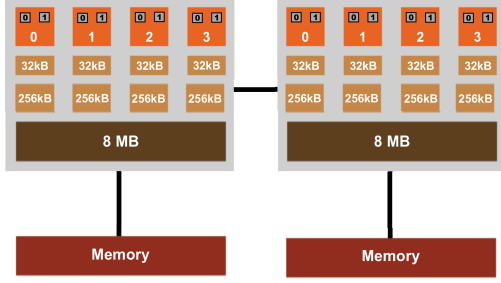


Figure 1: Thread and cache topology of an Intel Nehalem EP multicore dual-socket node




Figure 2: likwid-perfCtr: Interaction between event sets, hardware events and performance counters.

(e.g., 90% of all systems in the latest Top 500 list are of x86 type) we do not consider this a severe limitation. In other areas like, e.g., workstations or desktops, the x86 dominance is even larger.

In the following we describe the four tools in detail.

#### A. likwid-perfCtr

Hardware-specific optimization requires an intimate knowledge of the microarchitecture of a processor and the characteristics of the code. While many problems can be solved with profiling, common sense, and runtime measurements, additional information is often useful to get a complete picture.

Performance counters are facilities to count hardware events during code execution on a processor. Since this mechanism is implemented directly in hardware there is no overhead involved. All modern processors provide hardware performance counters, but their primary purpose is to support computer architects during the implementation phase. Still they are also attractive for application programmers, because they allow an in-depth view on what happens on the processor while running applications. There are generally two options for using hardware performance counter data: Either event counts are collected over the runtime of an application process (or probably restricted to certain code parts via an appropriate API), or overflowing hardware counters can generate interrupts, which can be used for IP or call-stack sampling. The latter option enables a very fine-grained view on a code’s resource requirements (limited only

by the inherent statistical errors). However, the first option is sufficient in many cases and also practically overhead-free. This is why it was chosen as the underlying principle for likwid-perfCtr.

The probably best known and widespread existing tool is the PAPI library [5], for which we provide a detailed comparison to likwid-perfCtr in Section III. A lot of research is targeted towards using performance counter data for automatic performance analysis and detecting potential performance bottlenecks [1], [2], [3]. However, those solutions are often too unwieldy for the common user, who would prefer a quick overview as a first step in performance analysis. A key design goal for likwid-perfCtr was ease of installation and use, minimal system requirements (no additional kernel modules and patches), and — at least for basic functionality — no changes to the user code. A prototype for the development of likwid-perfCtr is the SGI tool “perfex,” which was available on MIPS-based IRIX machines as part of the “SpeedShop” performance suite. Cray provides a similar, PAPI-based tool (craypat) on their systems. likwid-perfCtr offers comparable or improved functionality with regard to hardware performance counters on x86 processors, and is available as open source.

Hardware performance counters are controlled and accessed using processor-specific hardware registers (also called *model specific registers* (MSR)). likwid-perfCtr uses the Linux “msr” module to modify the MSRs from user space. The msr module is available in all Linux distributions with a 2.6 Linux kernel and implements the read/write access to MSRs based on device files.

likwid-perfCtr is a command line tool that can be used as a wrapper to an application. It allows simultaneous measurements on multiple cores. Events that are shared among the cores of a socket (this pertains to the “uncore” events on Core i7-type processors) are supported via “socket locks,” which enforce that all uncore event counts are assigned to one thread per socket. Events are specified on the command line, and the number of events to count concurrently is limited by the number of performance counters on the CPU. These features are available without any changes in the user’s source code. A small instrumentation (“marker”) API allows one to restrict measurements to certain parts of the code (named regions) with automatic accumulation over all regions of the same name. An important difference to most existing performance tools is that event counts are strictly core-based instead of process-based: Everything that runs and generates events on a core is taken into account; no attempt is made to filter events according to the process that caused them. The user is responsible for enforcing appropriate affinity to get sensible results. This could be achieved via likwid-pin (see below for more information):

```
$ likwid-perfCtr -c 1 \
-g SIMD_COMP_INST_RETIRED_PACKED_DOUBLE:PMC0,\
SIMD_COMP_INST_RETIRED_SCALAR_DOUBLE:PMC1 \
```

```
likwid-pin -c 1 ./a.out
```

(See below for typical output in a more elaborate setting.) In this example, the computational double precision packed and scalar SSE retired instruction counts on an Intel Core 2 processor are assigned to performance counters 0 and 1 and measured on core 1 over the duration of `a.out`'s runtime. The `likwid-pin` command is used here to bind the process to this core. As a side effect, it becomes possible to use `likwid-perfCtr` as a monitoring tool for a complete shared-memory node, just by specifying all cores for measurement and, e.g., “sleep” as an application:

```
$ likwid-perfCtr -c 0-7 \
  -g SIMD_COMP_INST_RETIRED_PACKED_DOUBLE:PMC0,\
    SIMD_COMP_INST_RETIRED_SCALAR_DOUBLE:PMC1 \
    sleep 1
```

Apart from naming events as they are documented in the vendor's manuals, it is also possible to use preconfigured *event sets* (groups) with derived metrics. This provides a simple abstraction layer in cases where standard information like memory bandwidth, FLOps per second, etc., is sufficient:

```
$ likwid-perfCtr -c 0-3 \
  -g FLOPS_DP ./a.out
```

At the time of writing, the following event sets are defined:

Event set	Function
FLOPS_DP	Double Precision MFlops/s
FLOPS_SP	Single Precision MFlops/s
L2	L2 cache bandwidth in MBytes/s
L3	L3 cache bandwidth in MBytes/s
MEM	Main memory bandwidth in MBytes/s
CACHE	L1 Data cache miss rate/ratio
L2CACHE	L2 Data cache miss rate/ratio
L3CACHE	L3 Data cache miss rate/ratio
DATA	Load to store ratio
BRANCH	Branch prediction miss rate/ratio
TLB	Translation lookaside buffer miss rate/ratio

The event groups are partly inspired from a technical report published by AMD [6]. We try to provide the same preconfigured event groups on all supported architectures, as long as the native events support them. This allows the beginner to concentrate on the useful information right away, without the need to look up events in the manuals (similar to PAPI's high-level events).

The interactions between event sets, hardware events, and performance counters are illustrated in Fig. 2. In the usage scenarios described so far there is no interference of `likwid-perfCtr` while user code is being executed, i.e., the overhead is very small (apart from the unavoidable API call overhead in marker mode). If the number of events is larger than the number of available counters, this mode of operation requires running the application more than once. For ease

of use in such situations, `likwid-perfCtr` also supports a *multiplexing mode*, where counters are assigned to several event sets in a “round robin” manner. On the downside, short-running measurements will then carry large statistical errors. Multiplexing is supported in wrapper and marker mode.

The following example illustrates the use of the marker API in a serial program with two named regions (“Main” and “Accum”):

```
#include <likwid.h>
...
int coreID = likwid_processGetProcessorId();
printf("Using_likwid\n");
likwid_markerInit(numberOfThreads,numberOfRegions);
int MainId = likwid_markerRegisterRegion("Main");
int AccumId = likwid_markerRegisterRegion("Accum");

likwid_markerStartRegion(0, coreID);
// measured code region
likwid_markerStopRegion(0, coreID, MainId);

for (j = 0; j < N; j++)
{
    likwid_markerStartRegion(0, coreID);
    // measured code region
    likwid_markerStopRegion(0, coreID, AccumId);
}

likwid_markerClose();
```

Event counts are automatically accumulated on multiple calls. Nesting or partial overlap of code regions is not allowed. The API requires specification of a thread ID (0 for one process only in the example) and the core ID of the thread/process. The `likwid` API provides simple functions to determine the core ID of processes or threads. The following listing shows the output of `likwid-perfCtr` after measurement of the FLOPS\_DP event group on four cores of an Intel Core 2 Quad processor in marker mode with two named regions (“Init” and “Benchmark,” respectively):

```
$ likwid-perfCtr -c 0-3 -g FLOPS_DP -m ./a.out
-----
CPU type:      Intel Core 2 45nm processor
CPU clock:     2.83 GHz
-----
Measuring group FLOPS_DP
-----
Region: Init
-----+-----+-----+-----+-----+
| Event | core 0 | core 1 | core 2 | core 3 |
+-----+-----+-----+-----+-----+
| INSTR_RETIRED_ANY | 313742 | 376154 | 355430 | 341988 |
| CPU_CLK_UNHALTED_CORE | 217578 | 504187 | 477785 | 459276 |
| SIMD_COMP_INST_RETIRED_PACKED_DOUBLE | 0 | 0 | 0 | 0 |
| SIMD_COMP_INST_RETIRED_SCALAR_DOUBLE | 1 | 1 | 1 | 1 |
+-----+-----+-----+-----+-----+
| Metric | core 0 | core 1 | core 2 | core 3 |
+-----+-----+-----+-----+-----+
| Runtime [s] | 7.67906e-05 | 0.000177945 | 0.000168626 | 0.000162094 |
| CPI | 0.693493 | 1.34037 | 1.34424 | 1.34296 |
| DP MFlops/s | 0.0130224 | 0.00561973 | 0.00593027 | 0.00616926 |
+-----+-----+-----+-----+-----+
Region: Benchmark
-----+-----+-----+-----+-----+
| Event | core 0 | core 1 | core 2 | core 3 |
+-----+-----+-----+-----+-----+
| INSTR_RETIRED_ANY | 1.88024e+07 | 1.85461e+07 | 1.84947e+07 | 1.84766e+07 |
| CPU_CLK_UNHALTED_CORE | 2.85838e+07 | 2.82369e+07 | 2.82429e+07 | 2.82066e+07 |
| SIMD...PACKED_DOUBLE | 8.192e+06 | 8.192e+06 | 8.192e+06 | 8.192e+06 |
| SIMD...SCALAR_DOUBLE | 1 | 1 | 1 | 1 |
+-----+-----+-----+-----+-----+
| Metric | core 0 | core 1 | core 2 | core 3 |
+-----+-----+-----+-----+-----+
```

Runtime [s]	0.0100882	0.00996574	0.00996787	0.00995505
CPI	1.52023	1.52252	1.52708	1.52661
DP MFlops/s	1624.08	1644.03	1643.68	1645.8

Note that the `INSTR_RETIRED_ANY` and `CPU_CLK_UNHALTED_CORE` events are always counted (using two unassignable “fixed counters” on the Core 2 architecture), so that the derived CPI metric (“cycles per instruction”) is easily obtained.

The following architectures are supported at the time of writing:

- Intel Pentium M (Banias, Dothan)
- Intel Atom
- Intel Core 2 (all variants)
- Intel Nehalem (all variants, including uncore events)
- Intel Westmere
- AMD K8 (all variants)
- AMD K10 (Barcelona, Shanghai, Istanbul)

### B. *likwid-topology*

Multicore/multisocket machines exhibit complex topologies, and this trend will continue with future architectures. Performance programming requires in-depth knowledge of cache and node topologies, e.g., about which caches are shared between which cores and which cores reside on which sockets. The Linux kernel numbers the usable cores and makes this information accessible in `/proc/cpuinfo`. Still how this numbering maps to the node topology depends on BIOS settings and may even differ for otherwise identical processors. The processor and cache topology can be queried with the `cpuid` machine instruction. `likwid-pin` is based directly on the data provided by `cpuid`. It extracts machine topology in an accessible way and can also report on cache characteristics. The thread topology is determined from the APIC (Advanced Programmable Interrupt Controller) ID. Starting with the Nehalem processor, Intel introduced a new `cpuid` leaf (0xB) to account for today’s more complex multicore chip topologies. Older Intel and AMD processors both have different methods to extract this information, all of which are supported by `likwid-topology`. Similar considerations apply for determining the cache topology. Starting with the Core 2 architecture Intel introduced the `cpuid` leaf 0x4 (deterministic cache parameters), which allows to extract the cache characteristics and topology in a systematic way. On older Intel processors the cache parameters were provided by means of a lookup table (`cpuid` leaf 0x2). AMD again has its own `cpuid` leaf for the cache parameters. The core functionality of `likwid-topology` is implemented in a C module, which can also be used as a library to access the information from within an application.

`likwid-topology` outputs the following information:

- Clock speed
- Thread topology (which hardware threads map to which physical resource)

- Cache topology (which hardware threads share a cache level)
- Extended cache parameters for data caches.

The following output was obtained on an Intel Nehalem EP Westmere processor and includes extended cache information:

```
$ likwid-topology -c
-----
CPU name:      Unknown Intel Processor
CPU clock:     2.93 GHz

*****
Hardware Thread Topology
*****
Sockets:      2
Cores per socket: 6
Threads per core: 2
-----
HWThread      Thread      Core      Socket
0              0              0          0
1              0              1          0
2              0              2          0
3              0              8          0
4              0              9          0
5              0             10          0
6              0             0          1
7              0              1          1
8              0              2          1
9              0              8          1
10             0              9          1
11             0             10          1
12             1              0          0
13             1              1          0
14             1              2          0
15             1              8          0
16             1              9          0
17             1             10          0
18             1              0          1
19             1              1          1
20             1              2          1
21             1              8          1
22             1              9          1
23             1             10          1
-----
Socket 0: ( 0 12 1 13 2 14 3 15 4 16 5 17 )
Socket 1: ( 6 18 7 19 8 20 9 21 10 22 11 23 )
-----

*****
Cache Topology
*****
Level:  1
Size:   32 kB
Type:   Data cache
Associativity: 8
Number of sets: 64
Cache line size: 64
Inclusive cache
Shared among 2 threads
Cache groups: ( 0 12 ) ( 1 13 ) ( 2 14 ) ( 3 15 ) ( 4 16 )
              ( 5 17 ) ( 6 18 ) ( 7 19 ) ( 8 20 ) ( 9 21 ) ( 10 22 ) ( 11 23 )
-----
Level:  2
Size:   256 kB
Type:   Unified cache
Associativity: 8
Number of sets: 512
Cache line size: 64
Inclusive cache
Shared among 2 threads
Cache groups: ( 0 12 ) ( 1 13 ) ( 2 14 ) ( 3 15 ) ( 4 16 )
              ( 5 17 ) ( 6 18 ) ( 7 19 ) ( 8 20 ) ( 9 21 ) ( 10 22 ) ( 11 23 )
-----
Level:  3
Size:   12 MB
Type:   Unified cache
Associativity: 16
Number of sets: 12288
Cache line size: 64
Non Inclusive cache
Shared among 12 threads
Cache groups: ( 0 12 1 13 2 14 3 15 4 16 5 17 )
              ( 6 18 7 19 8 20 9 21 10 22 11 23 )
-----
```

One can also get an accessible overview of the node’s cache and socket topology in ASCII art (via the `-g` option). The following listing fragment shows the output for the same chip as above. Note that only one socket is shown (belonging to the first L3 cache group above):

```
+-----+
| 0 12 | | 1 13 | | 2 14 | | 3 15 | | 4 16 | | 5 17 | |
+-----+
```







Figure 3: Basic structure of likwid-pin.



### C. likwid-pin

Thread/process affinity is vital for performance. If topology information is available, it is possible to pin threads according to the application’s resource requirements like bandwidth, cache sizes, etc. Correct pinning is even more important on processors supporting SMT, where multiple hardware threads share resources on a single core. likwid-pin supports thread affinity for all threading models that are based on POSIX threads, which includes most OpenMP implementations. By overloading the `pthread_create` API call with a shared library wrapper, each thread can be pinned in turn upon creation, working through a list of core IDs. This list, and possibly other parameters, are encoded in environment variables that are evaluated when the library wrapper is first called. likwid-pin simply starts the user application with the library preloaded.

The overall mechanism is illustrated in Fig. 3. No code changes are required, but the application must be dynamically linked. This mechanism is independent of processor architecture, but the way the compiled code creates application threads must be taken into account: For instance, the Intel OpenMP implementation always runs `OMP_NUM_THREADS+1` threads but uses the first newly created thread as a management thread, which should not be pinned. This knowledge must be conveyed to the wrapper library. The following example shows how to use likwid-pin with an OpenMP application compiled with the Intel compiler:

```
$ export OMP_NUM_THREADS=4
$ likwid-pin -c 0-3 -t intel ./a.out
```

Currently, POSIX threads, Intel OpenMP, and GNU (gcc)

OpenMP are supported, and the latter is assumed as the default if no `-t` switch is given. Other threading implementations are supported via a “skip mask.” This mask is interpreted as a binary pattern and specifies which threads should not be pinned by the wrapper library (the explicit mask for Intel binaries would be `0x1`). The skip mask makes it possible to pin hybrid applications as well by skipping MPI shepherd threads. For Intel-compiled binaries using the Intel MPI library, the appropriate skip mask is `0x3`:

```
$ export OMP_NUM_THREADS=8
$ mpiexec -n 64 -pernode \
    likwid-pin -c 0-7 -s 0x3 ./a.out
```

This would start 64 MPI processes on 64 nodes (via the `-pernode` option) with eight threads each, and not bind the first two newly created threads.

In general, likwid-pin can be used as a replacement for the `taskset` tool, which cannot pin threads individually. Note, however, that likwid-pin, in contrast to `taskset`, does not establish a Linux `cputset` in which to run the application.

Some compilers have their own means for enforcing thread affinity. In order to avoid interference effects, those mechanisms should be disabled when using likwid-pin. In case of recent Intel compilers, this can be achieved by setting the environment variable `KMP_AFFINITY` to `disabled`. The current version of LIKWID does this automatically.

The big advantage of likwid-pin is its portable approach to the pinning problem, since the same tool can be used for all applications, compilers, MPI implementations, and processor types. In Section IV-A the usage model is analyzed in more detail on the example of the STREAM triad.

### D. likwid-features

An important hardware optimization on modern processors is to hide data access latencies by hardware prefetching. Intel processors not only have a prefetcher for main memory; several prefetchers are responsible for moving data between cache levels. Often it is beneficial to know the influence of the hardware prefetchers. In some situations turning off hardware prefetching even increases performance. On the Intel Core 2 processor this can be achieved by setting bits in the `IA32_MISC_ENABLE` MSR register. likwid-features allows viewing and altering the state of these bits. Besides the ability to toggle the hardware prefetchers, likwid-features also reports on the state of switchable processor features like, e.g., Intel Speedstep:

```
$ likwid-features
-----
CPU name:      Intel Core 2 65nm processor
CPU core id:   0
-----
Fast-Strings:      enabled
Automatic Thermal Control:  enabled
Performance monitoring:  enabled
Hardware Prefetcher:  enabled
Branch Trace Storage:  supported
PEBS:              supported
Intel Enhanced SpeedStep:  enabled
MONITOR/MWAIT:      supported
Adjacent Cache Line Prefetch:  enabled
Limit CPUID Maxval:  disabled
```

```

XD Bit Disable:      enabled
DCU Prefetcher:     enabled
Intel Dynamic Acceleration: disabled
IP Prefetcher:      enabled

```

Disabling, e.g., adjacent cache line prefetch then works as follows:

```

$ likwid-features -u CL_PREFETCHER
[...]
CL_PREFETCHER: disabled

```

likwid-features currently only works for Intel Core 2 processors, but support for other architectures is planned for the future.

### III. COMPARISON WITH PAPI

PAPI [4] is a popular and well known framework to measure hardware performance counter data. In contrast to LIKWID it relies on other software to implement the architecture-specific parts and concentrates on providing a portable interface to performance metrics on various platforms and architectures. PAPI is mainly intended to be used as a library but also includes a small collection of command line utilities. At the time of writing PAPI is available in a classic version (PAPI 3.7.2) and a new main branch (PAPI 4.0.0). Both version rely on autoconf to generate the build configuration.

Table I compares PAPI with LIKWID without any claim for completeness. Of course many issues are difficult to quantify, and a thorough coverage of these points is beyond the scope of this paper. Still the comparison should give an impression about the differences between both tools. The most important difference is that LIKWID main focus is in providing a collection of command line tools for the end user while PAPI's main focus is to be used as a library by other tools.

### IV. CASE STUDIES

#### A. Case study 1: Influence of thread topology on STREAM triad performance

To illustrate the general importance of thread affinity we use the well known OpenMP STREAM triad on an Intel Westmere dual-socket system. Intel Westmere is a hexacore design based on the Nehalem architecture and supports two SMT threads per physical core. Two different compilers are considered: Intel icc (11.1, with options `-openmp -O3 -xSSE4.2 -fno-fnalias`) and gcc (4.3.3, with options `-O3 -fopenmp -fargument-noalias`). The executable for the test on AMD Istanbul was compiled with Intel icc (11.1, `-openmp -O3 -fno-fnalias`). Intel compilers support thread affinity only if the application is executed on Intel processors. The functionality of this topology interface is controlled by setting the environment variable `KMP_AFFINITY`. In our tests `KMP_AFFINITY` was set to `disabled`. For the case of the STREAM triad




Figure 4: STREAM triad test run for the Intel icc compiler on a two-socket 12-core Westmere system with 100 samples per thread count (this will be the same for all subsequent test runs). The application is not explicitly pinned. The box plot shows the 25-50 range with the median line.




Figure 5: STREAM triad test run for the Intel icc compiler. The application is pinned such that threads are equally distributed on the sockets to utilize the memory bandwidth in the most effective way. Moreover the threads are first distributed over physical cores and then over SMT threads.




Figure 6: STREAM triad test run for the Intel icc compiler. The application was run with the affinity interface of the Intel OpenMP implementation set to "scatter."

	LIKWID	PAPI
Dependencies	Needs system headers of Linux 2.6 kernel. No other external dependencies.	Needs kernel patches depending on platform and architecture. No patches necessary on Linux kernels > 2.6.31.
Installation	Build system based on make only. Install documentation 10 lines. Build configuration in a single text file (21 lines).	Install documentation is 582 lines (3.7.2) and 397 lines (4.0.0). The installation of PAPI for this comparison was not without problems.
Command line tools	Core is a collection of command line tools which are intended to be used standalone.	Collection of small utilities. These utilities are not supposed to be used as standalone tools. There are many PAPI-based tools available from other sources.
User API support	Simple API for configuring named code regions. API only turns counters on and off. Configuration of events and output of results is still based on the command line tool.	Comparatively high-level API. Events must be configured in the code.
Library support	While it can be used as library this was not initially intended.	Mature and well tested library API for building own tooling.
Topology information	Listing of thread and cache topology. Results are extracted from cpuid and presented in an accessible way as text and ASCII art. Nondata caches are omitted. No output of TLB information.	Information also based on cpuid. Utility outputs all caches (including TLBs). No output of shared cache information. Thread topology only as accumulated counts of HW threads and Cores. No mapping from processor Ids to thread topology.
Thread and process pinning	There is a dedicated tool for pinning processes and threads in a portable and simple manner. This tool is intended to be used together with likwid-perfCtr	No support for pinning.
Multicore support	Multiple cores can be measured simultaneously. Binding of threads or processes to correct cores is the responsibility of the user.	No explicit support for multicore measurements.
Uncore support	Uncore events are handled by applying socket locks, which prevent multiple measurements in threaded mode.	No explicit support for measuring shared resources.
Event abstraction	Preconfigured event sets (so-called event groups) with derived metrics.	Abstraction through papi_events, which map to native events.
Platform support	Supports only x86-based processors on Linux with 2.6 kernel.	Supports a wide range of architectures on various platforms (dedicated support for HPC systems like BlueGene or Cray XT3/4/5) with various operating systems (Linux, FreeBSD, and Windows).
Correlated measurements	LIKWID can measure performance counters only	PAPI-C can be extended to measure and correlate various data like, e.g., fan speeds or temperatures.

Table I: Comparison between LIKWID and PAPI

on these ccNUMA architectures the best performance is achieved if threads are equally distributed across the two sockets.

Figure 4 shows the results for the Intel compiler with no explicit pinning. In contrast, the data in Fig. 5 was obtained with the threads distributed in a round-robin manner across physical sockets using likwid-pin. As described earlier, the Intel OpenMP implementation creates `OMP_NUM_THREADS` in addition to the initial master thread, but the first newly created thread is used as a “shepherd” and must not be pinned. likwid-pin provides a type parameter to indicate the OpenMP implementation and automatically sets an appropriate skip mask. In contrast, gcc OpenMP only creates `OMP_NUM_THREADS-1` additional threads and does not require a shepherd thread. As can be seen in Fig. 4, the non-pinned runs show a large variance in performance especially for the smaller thread counts where the probability is large that only one socket is used. With larger thread counts there is a high probability that both sockets are used, still there is also a chance that cores are

oversubscribed and performance is thereby reduced. The pinned case consistently shows high performance.

The effectiveness of the affinity functionality of the Intel OpenMP implementation can be seen in Fig. 6. This option provides the same high performance as with likwid-pin, at all thread counts.

In Fig. 7 and Fig. 8 the same test is shown for gcc. Interestingly, the performance distribution is significantly different compared to the non-pinned Intel icc test case in Fig. 4. While with Intel icc the variance was larger for smaller thread counts, for gcc the variance for this region is small and results are bad with high probability. For larger thread counts this picture is reversed: Intel icc has a small variance while gcc shows the biggest variance. One possible explanation is that the gcc code is less dense in terms of cycles per instruction, tolerating an oversubscription, and can probably benefit from SMT threads to a larger extent than the Intel icc code. This behavior was not investigated in more detail here.

Finally the Intel icc executable was also benchmarked on

a two-socket AMD Istanbul hexacore node. Fig. 9 shows that there is a large performance variance in the unpinned case, as expected. Still no significant difference can be seen between the distribution for smaller or larger thread counts. Enforcing affinity with likwid-pin (Fig. 10) yields good, stable results for all thread counts. It is apparent that the SMT threads of Intel Westmere increase the probability for interference of competing processes. It also makes Intel Westmere more sensitive to oversubscription and leads to volatile performance with smaller thread counts.




Figure 7: STREAM triad test run for the gcc compiler without pinning.




Figure 8: STREAM triad test run for the gcc compiler. The application was pinned with likwid-pin. The arguments for likwid-pin and the plot properties are the same as for the Intel icc test in Fig. 5.

### B. Case Study 2: Influence of thread topology on a topology-aware stencil code

While in the first case study the ccNUMA characteristics of the benchmark systems only required the distribution of threads across cores to be “uniform,” the following example will show that the specific thread and cache topology must sometimes be taken into account, and the exact mapping of threads to cores becomes vital for getting good performance.

We investigated a highly optimized application that was specifically designed to utilize the shared caches of modern multicore architectures. It implements an iterative 3D




Figure 9: STREAM triad test run for the Intel icc compiler on an AMD Istanbul node without pinning.




Figure 10: STREAM triad test run for the Intel icc compiler on a AMD Istanbul node. The application was pinned with likwid-pin. The arguments for likwid-pin are the same as in Fig. 5.

Jacobi smoother using a 7-point stencil and is based on the POSIX threads library. All critical computational kernels are implemented in assembly language. This code uses implicit temporal blocking based on a pipeline parallel processing approach [8]. The benchmarks were performed on a dual-socket Intel Nehalem EP quad-core system. Figure 11 shows that in case of wrong pinning the effect of the optimization is reversed and performance is reduced by a factor of two, because the shared cache cannot be leveraged to increase the computational intensity. In this case performance is even lower than with a naive threaded baseline code without temporal blocking. Hence, just pinning threads “evenly” through the machine is not sufficient here; the topology of the machine requires a very specific thread-core mapping for the blocking optimizations to become effective.

### C. Case Study 3: Examining the effect of temporal blocking

Using the code from the case study in Sec. IV-B, we performed hardware performance counter measurements with likwid-perfCtr on a dual-socket Nehalem EP system to quantify the effect of a temporal blocking optimization. The measurements use three versions of a 7-point stencil Jacobi






Figure 11: Performance of an optimized 3D Jacobi smoother versus linear problem size (cubic computational grid) on a dual-socket Intel Nehalem EP node (2.66 GHz) using one thread group consisting of four threads, pinned to the physical cores of one socket (circles). In contrast, pinning pairs of threads to different sockets (squares) is hazardous for performance. The threaded baseline with nontemporal stores is shown for reference (triangles). Results are in million lattice site updates per second [MLUPS].

kernel: (i) a standard threaded code with temporal stores (“threaded”), (ii) the same threaded implementation with nontemporal stores (“threaded (NT)”), and (iii) the temporal blocking code mentioned in the previous section. The data transfer volume to and from main memory is used as a metric to evaluate the effect of temporal blocking. Two uncore events are relevant here: The number of cache lines allocated in L3, and the number of cache lines victimized from L3 (see Tab.II). In all cases, the same number of stencil updates was executed with identical settings, and the four physical cores of one socket were utilized. The results are shown in Table II. It can be seen that nontemporal stores save about 1/3 of the data transfer volume compared to the code with temporal stores, because the write allocate on store misses is eliminated. The optimized version again reduces the data transfer volume significantly, as expected. However, the 4.5-fold overall decrease in memory traffic does not translate into a proportional performance boost. There are two reasons for this failure: (i) One data stream towards main memory cannot fully utilize the memory bandwidth on the Nehalem EP, while the standard threaded versions are able to saturate the bus. (ii) The performance difference between the saturated main memory case and the L3 bandwidth for Jacobi is small (compared to other, more bandwidth-starved designs), which limits the performance benefit of temporal blocking on this processor. See [9] for a performance model that describes those effects.

## V. CONCLUSION AND FUTURE PLANS

LIKWID is a collection of command line applications supporting performance-oriented software developers in their effort to utilize today’s multicore processors in an effective manner. LIKWID does not try to follow the trend to provide yet another complex and sophisticated tooling

environment, which would be difficult to set up and would overwhelm the average user with large amounts of data. Instead it tries to make the important functionality accessible with as few obstacles as possible. The focus is put on simplicity and low overhead. likwid-topology and likwid-pin enable the user to account for the influence of thread and cache topology on performance and pin their application to physical resources in all possible scenarios with one single tool and no code changes. Prototypically we have shown the influence of thread topology and correct pinning on the example of the STREAM triad benchmark. Moreover thread pinning and performance characteristics were reviewed for an optimized topology-aware stencil code using likwid-perfCtr. LIKWID is open source and released under GPL2. It can be downloaded at <http://code.google.com/p/likwid/>.

LIKWID is still in alpha stage. Near-term goals are to consolidate the current features and release a stable version, and to include support for more processor types. An important feature missing in likwid-topology is to include NUMA information in the output. likwid-pin will be equipped with cpuset support, so that logical core IDs may be used when binding threads. Further goals are the combination of LIKWID with one of the available MPI profiling frameworks to facilitate the collection of performance counter data in MPI programs. Most of these frameworks rely on the PAPI library at the moment.

Future plans include applying the philosophy of LIKWID to other areas like, e.g., profiling (also on the assembly level) and low-level benchmarking with a tool creating a “bandwidth map.” This will allow a quick overview of the cache and memory bandwidth bottlenecks in a shared-memory node, including the ccNUMA behavior. It is also planned to port parts of LIKWID to the Windows operating system. On popular demand, future releases will also include support for XML output.

## ACKNOWLEDGMENT

We are indebted to Intel Germany for providing test systems and early access hardware for benchmarking. Many thanks to Michael Meier, who had the basic idea for likwid-pin, implemented the prototype, and provided many useful thoughts in discussions. This work was supported by the Competence Network for Scientific and Technical High Performance Computing in Bavaria (KONWIHR) under the project “OMI4papps.”

	threaded	threaded (NT)	blocked
UNC_L3_LINES_IN_ANY	$5.91 \cdot 10^8$	$3.44 \cdot 10^8$	$1.30 \cdot 10^8$
UNC_L3_LINES_OUT_ANY	$5.87 \cdot 10^8$	$3.43 \cdot 10^8$	$1.29 \cdot 10^8$
Total data volume [GB]	75.39	43.97	16.57
Performance [MLUPS]	784	1032	1331

Table II: likwid-perfCtr measurements on one Nehalem EP socket, comparing the standard threaded Jacobi solver with and without nontemporal stores with a temporally blocked variant.

# REFERENCES

- [1] G. Jost, J. Haoqiang, J. Labarta, J. Gimenez, J. Caubet: *Performance analysis of multilevel parallel applications on shared memory architectures*. Proceedings of the Parallel and Distributed Processing Symposium, 2003.
- [2] M. Gerndt, E. Kereku: *Automatic Memory Access Analysis with Periscope*. ICCS '07: Proceedings of the 7th international conference on Computational Science, 2007, 847–854.
- [3] M. Gerndt, K. F rlinger, E. Kereku: *Periscope: Advanced Techniques for Performance Analysis*. PARCO, 2005, 15–26.
- [4] D. Terpstra, H. Jagode, H. You, J. Dongarra: *Collecting Performance Data with PAPI-C*. Proceedings of the 3rd Parallel Tools Workshop, Springer Verlag, Dresden, Germany, 2010.
- [5] S. Browne, C. Deane, G. Ho, P. Mucci: *PAPI: A Portable Interface to Hardware Performance Counters*. Proceedings of Department of Defense HPCMP Users Group Conference, June 1999.
- [6] Paul J. Drongowski: *Basic Performance Measurements for AMD Athlon 64, AMD Opteron and AMD Phenom Processors*. Technical Note, Advanced Micro Devices, Inc. Boston Design Center, September 2008.
- [7] L. DeRose, B. Homer, and D. Johnson: *Detecting application load imbalance on high end massively parallel systems*. Euro-Par 2007 Parallel Processing Conference, 2007, 150–159.
- [8] J. Treibig, G. Wellein, G. Hager: *Efficient multicore-aware parallelization strategies for iterative stencil computations*. Submitted to Journal of Computational Science, 2010. Preprint <http://arxiv.org/abs/1004.1741>.
- [9] M. Wittmann, G. Hager and G. Wellein. *Multicore-aware parallel temporal blocking of stencil codes for shared and distributed memory*. Workshop on Large-Scale Parallel Processing 2010 (IPDPS2010), Atlanta, GA, April 23, 2010. Preprint <http://arxiv.org/abs/0912.4506>